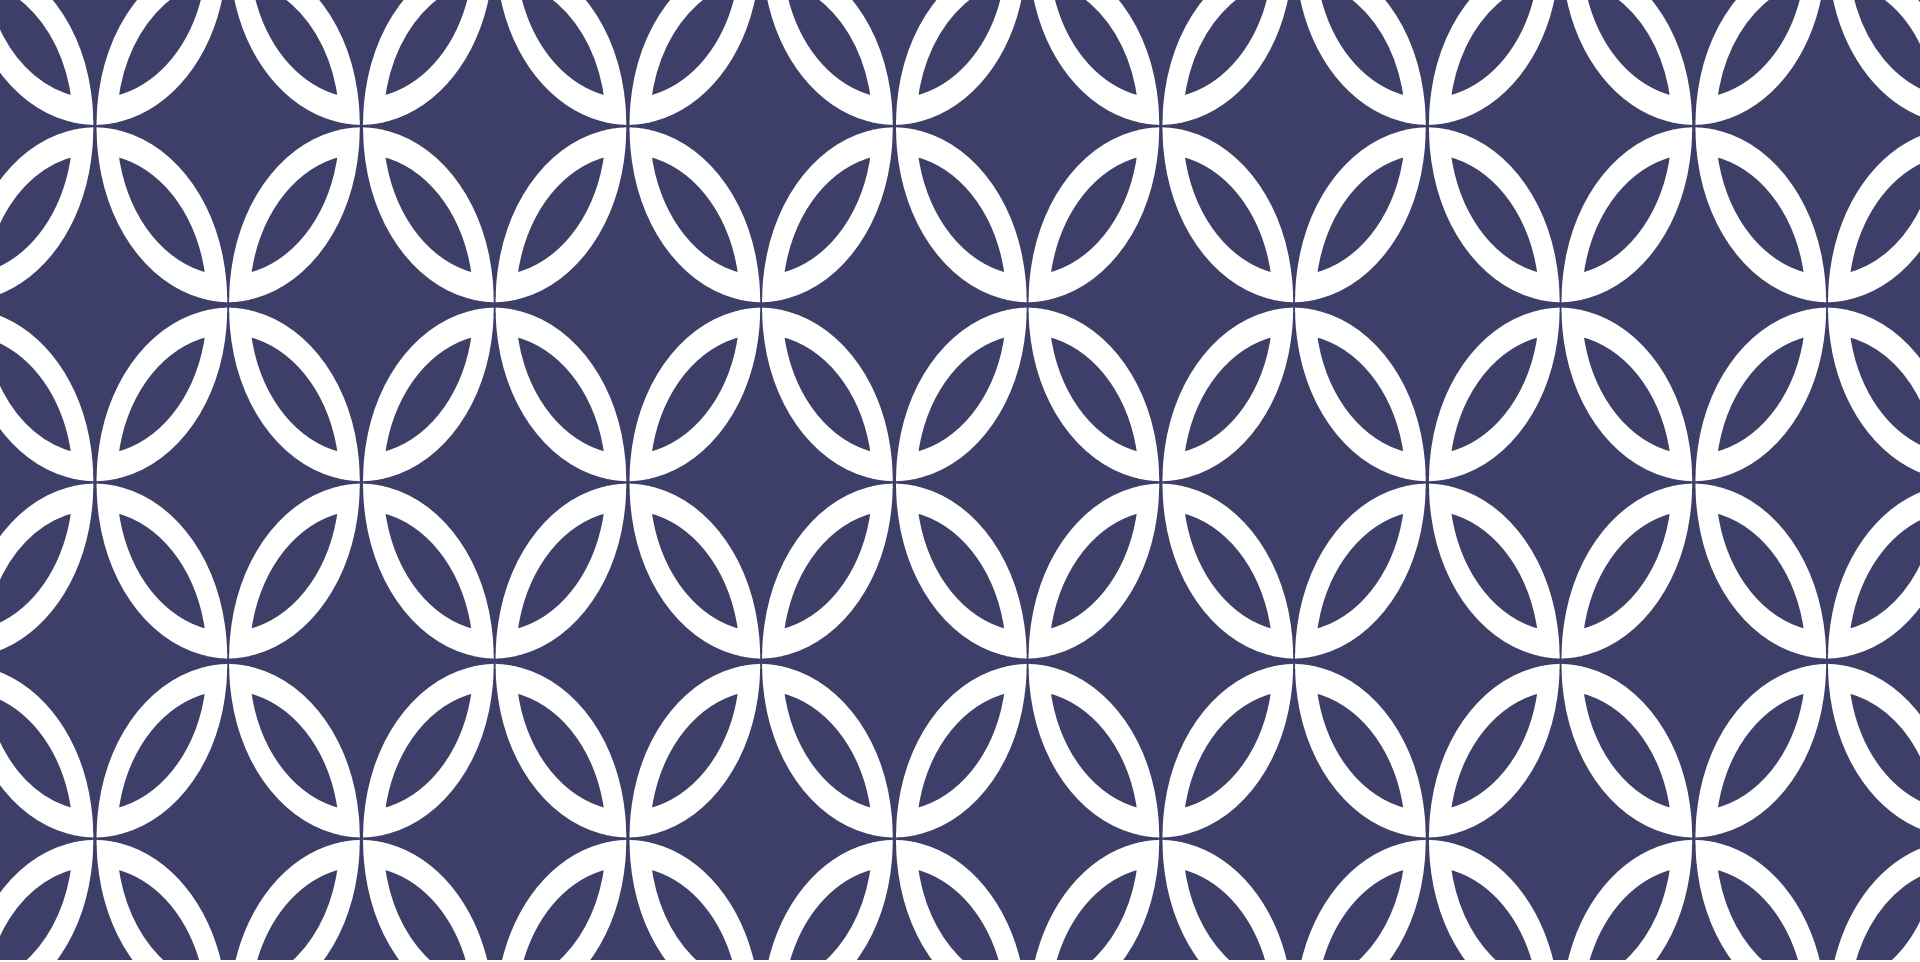


ACADGILD

Presents

Introduction to Big Data and Hadoop





Session 2 – Hadoop Framework Description



Agenda – Hadoop


1. Solving Big Data Problem
2. Hadoop Cluster
 - Introduction
 - Concepts
 - Why Hadoop Cluster?
 - Why Hadoop2 Came About?
 - How Hadoop Works
3. Hadoop 1.x Architecture
4. Progression from Hadoop 1.x to Hadoop-2.x
 - Core Components of Hadoop
 - NameNode Backup in Hadoop1.x
 - HDFS – High Availability Feature in Hadoop 2.x
5. Introduction to a YARN Application
6. Anatomy of a YARN Application
 - Phase I
 - Phase II



Breaking Big Data into Chunks ...

Can you process a file of 1 TB size on your desktop or laptop ?

Nooooooo !!





Hadoop Cluster Concepts

Hadoop cluster works as follows:

- Data is broken into chunks & distributed to different nodes.
- Each node(computer) processes data in it's storage to produce local results
- Results from each node are assembled to produce final results.



Why Hadoop Cluster?

- Data Storage has grown exponentially in recent past but data reading speed has not improved radically.

Data Read Speed Comparisons over time

Year	Data Size	Transfer Speed	Time Taken
1990	1400 MB	4.5 MB/s	5 Minutes
2010	1 TB	100 MB/s	3 Hours
Hadoop Results			
2013	1TB	100 Drives	2 Minutes




Why Hadoop 2 Came About?

MapReduce (MRv1) has following issues:

- Scalability Issues (max of 4000 nodes only) in MapReduce1
- Inflexible Resource Management Issue
 - MapReduce1 had slot based model.
 - Each TaskTracker is configured to have N slots at start-up.
 - A task is executed in a single slot.
 - Slots are configured on maximum memory at start-up (if more memory is freed up later on, it is unutilized).

How Hadoop Works





Hadoop1.x Architecture

Hadoop1.0 Architecture:

MapReduce

Job Tracker

Master Node

Task Tracker

MapReduce

Slave Node

Task Tracker

MapReduce

Slave Node


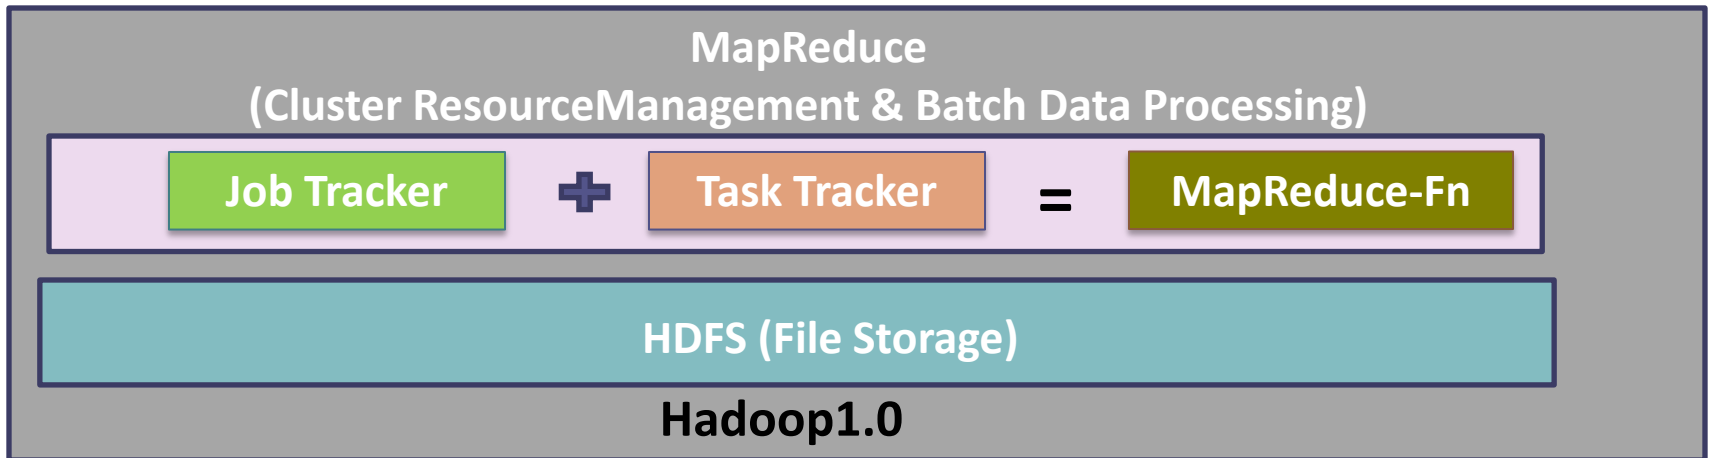
MapReduce
(Cluster ResourceManagement
& Batch Data Processing)

HDFS (File
Storage)

Hadoop1.0


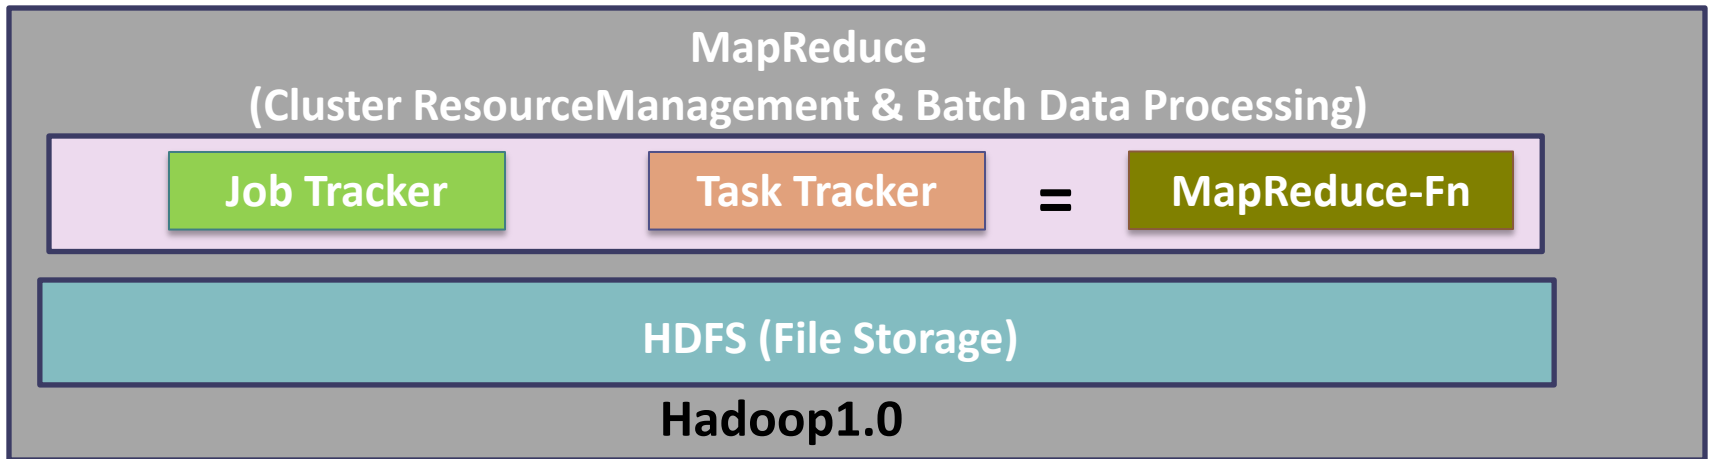


Progression from Hadoop1.x to Hadoop2




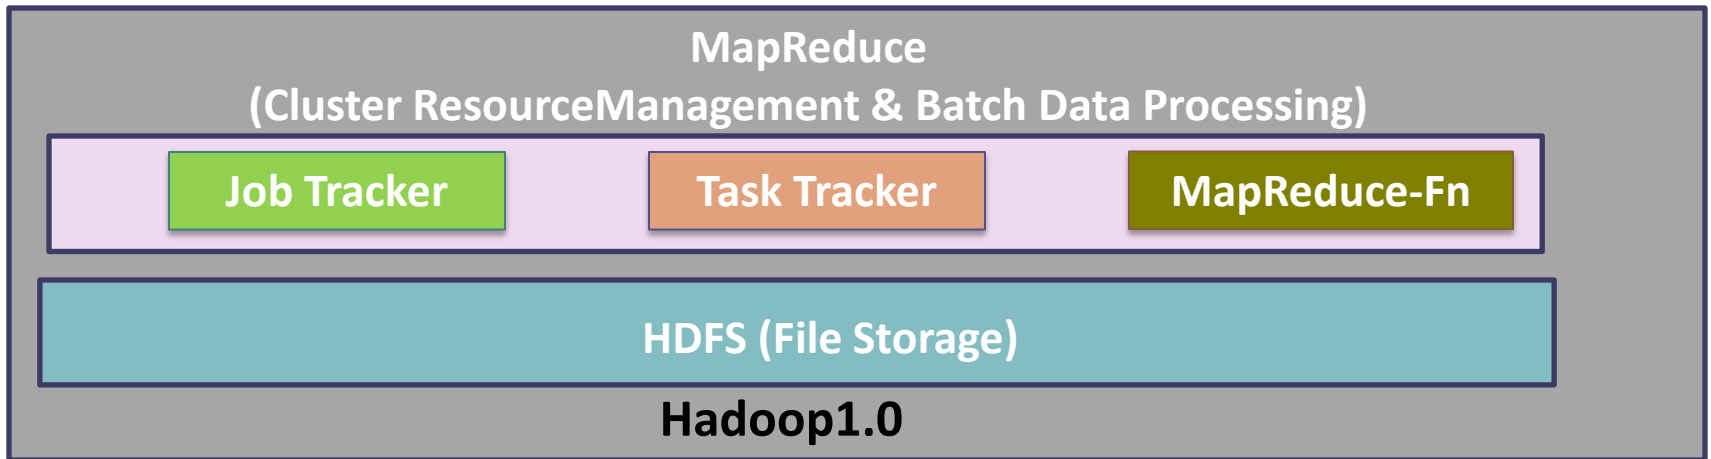


Progression from Hadoop1.x to Hadoop2




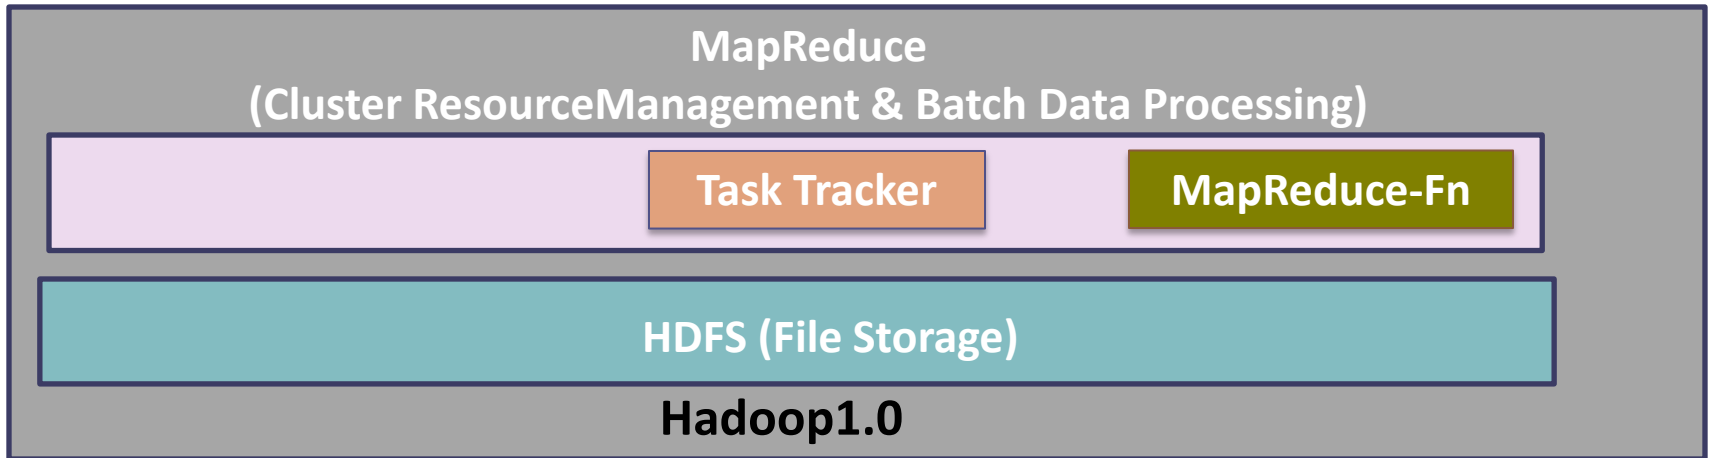


Progression from Hadoop1.x to Hadoop2




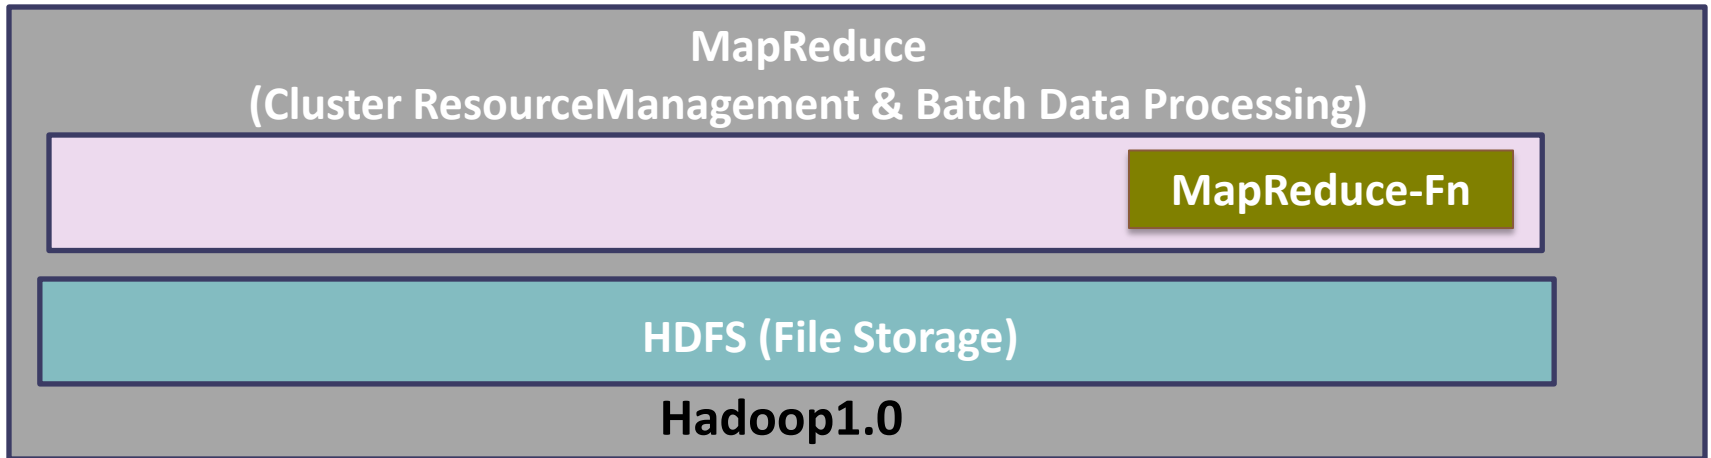


Progression from Hadoop1.x to Hadoop2



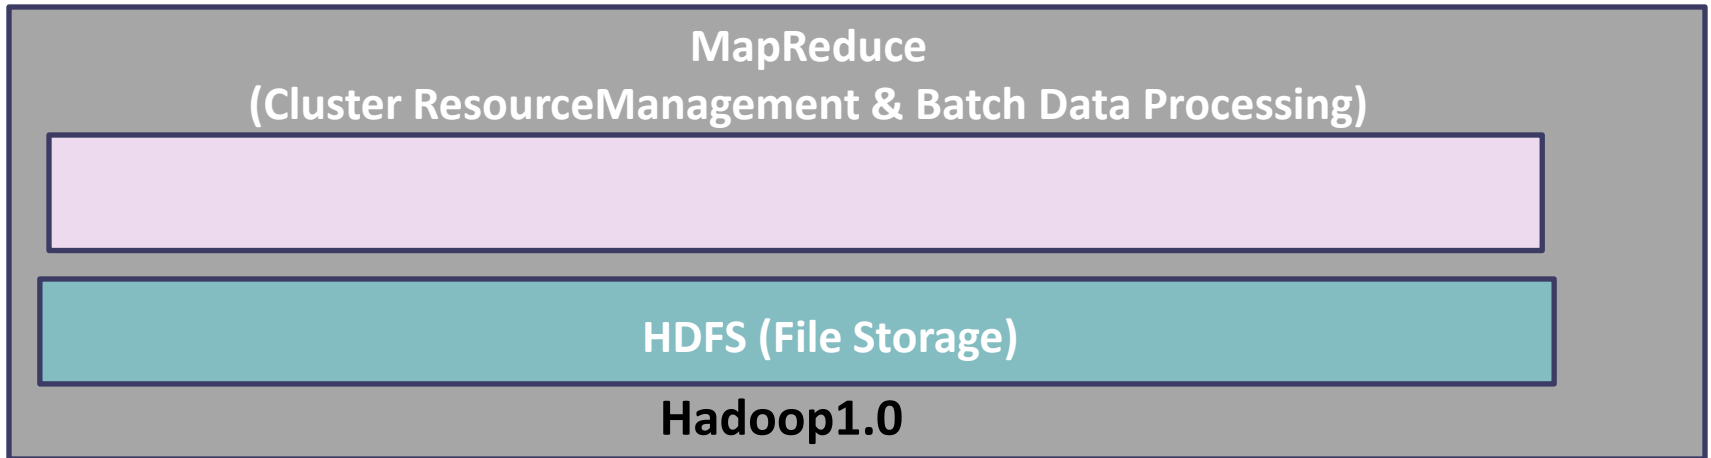


Progression from Hadoop1.x to Hadoop2






Progression from Hadoop1.x to Hadoop2

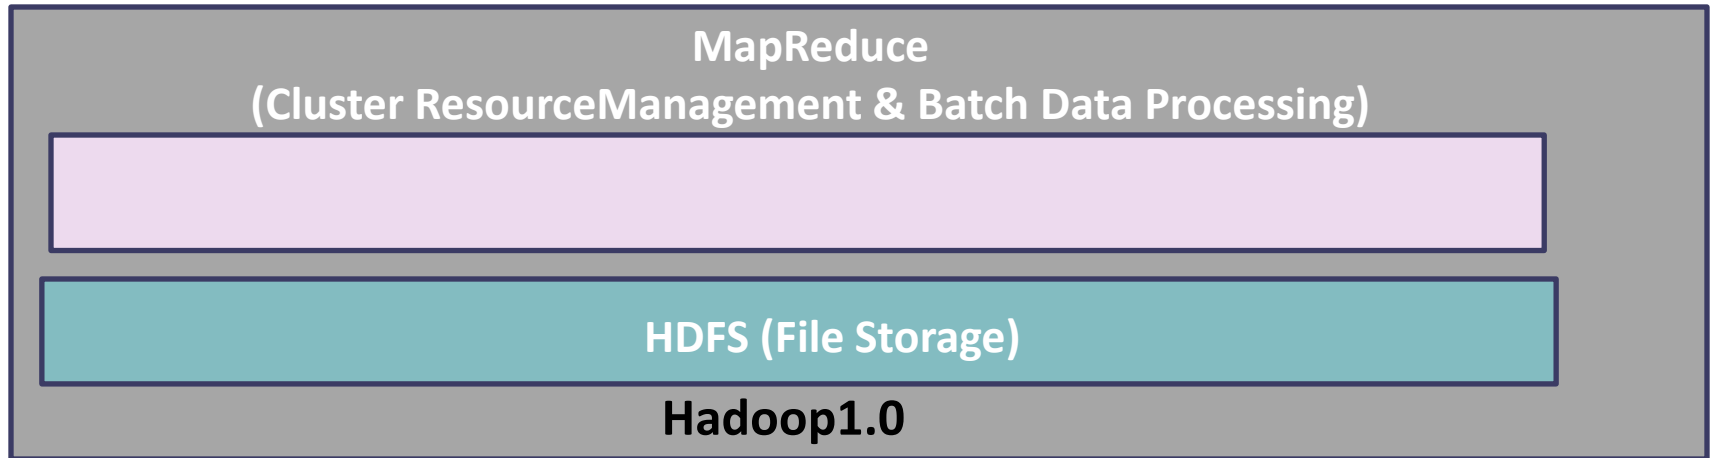


MapReduce-Fn






Progression from Hadoop1.x to Hadoop2



MapReduce-Fn





Core Components of Hadoop Cluster



Core Components of Hadoop Cluster

Hadoop 2.x Core Components



Core Components of Hadoop Cluster

Hadoop 2.x Core Components

Storage





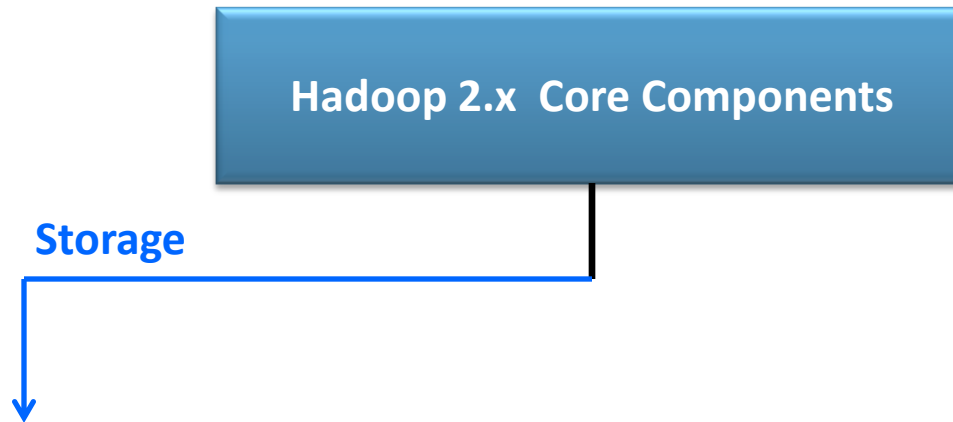
Core Components of Hadoop Cluster

Hadoop 2.x Core Components

Storage

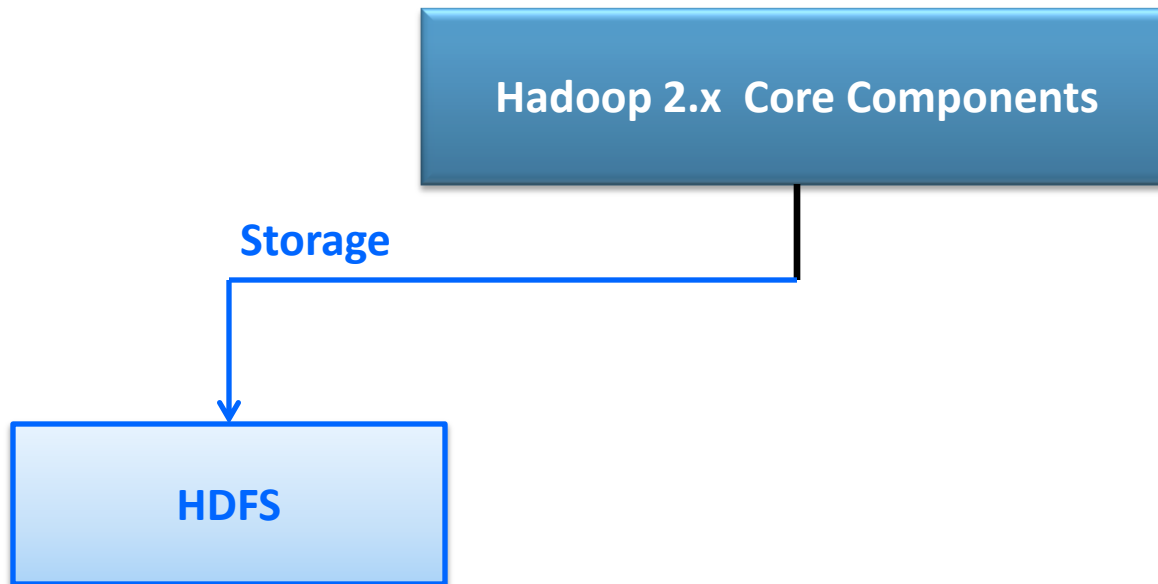


Core Components of Hadoop Cluster



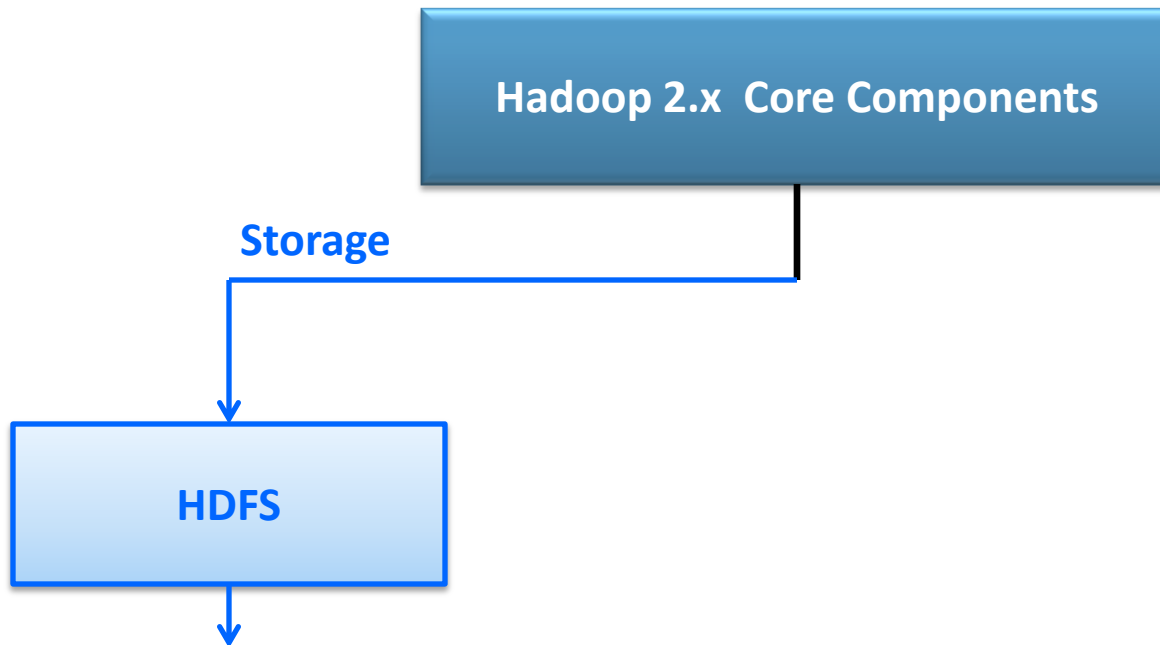


Core Components of Hadoop Cluster



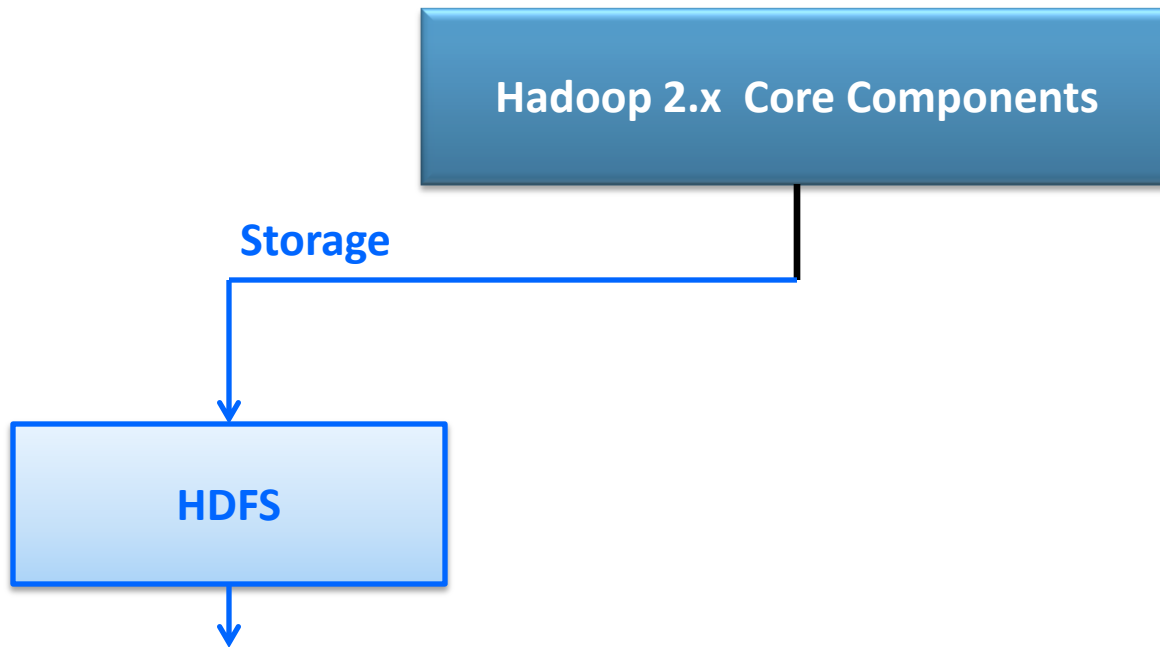


Core Components of Hadoop Cluster



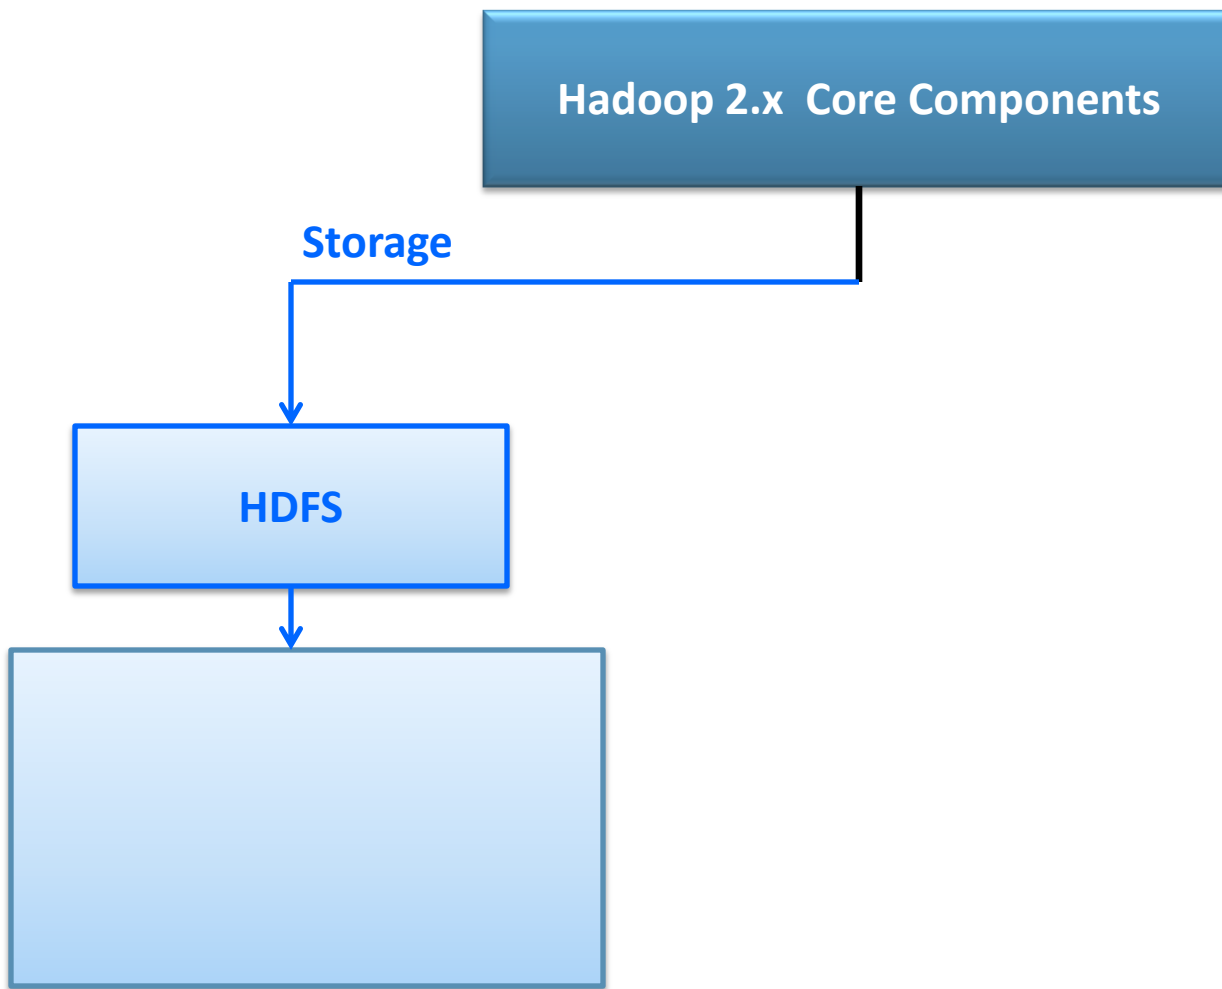


Core Components of Hadoop Cluster



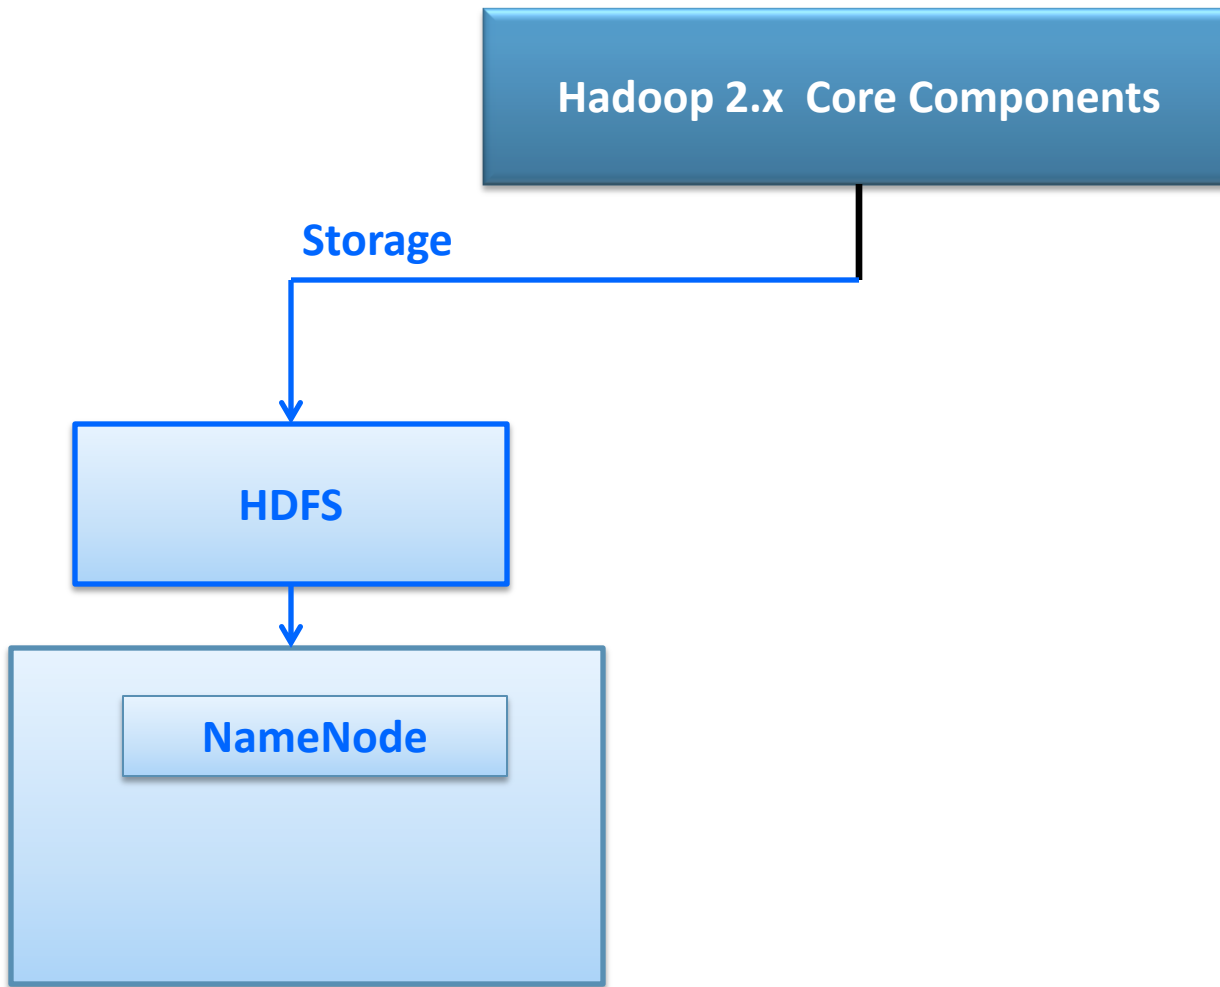


Core Components of Hadoop Cluster



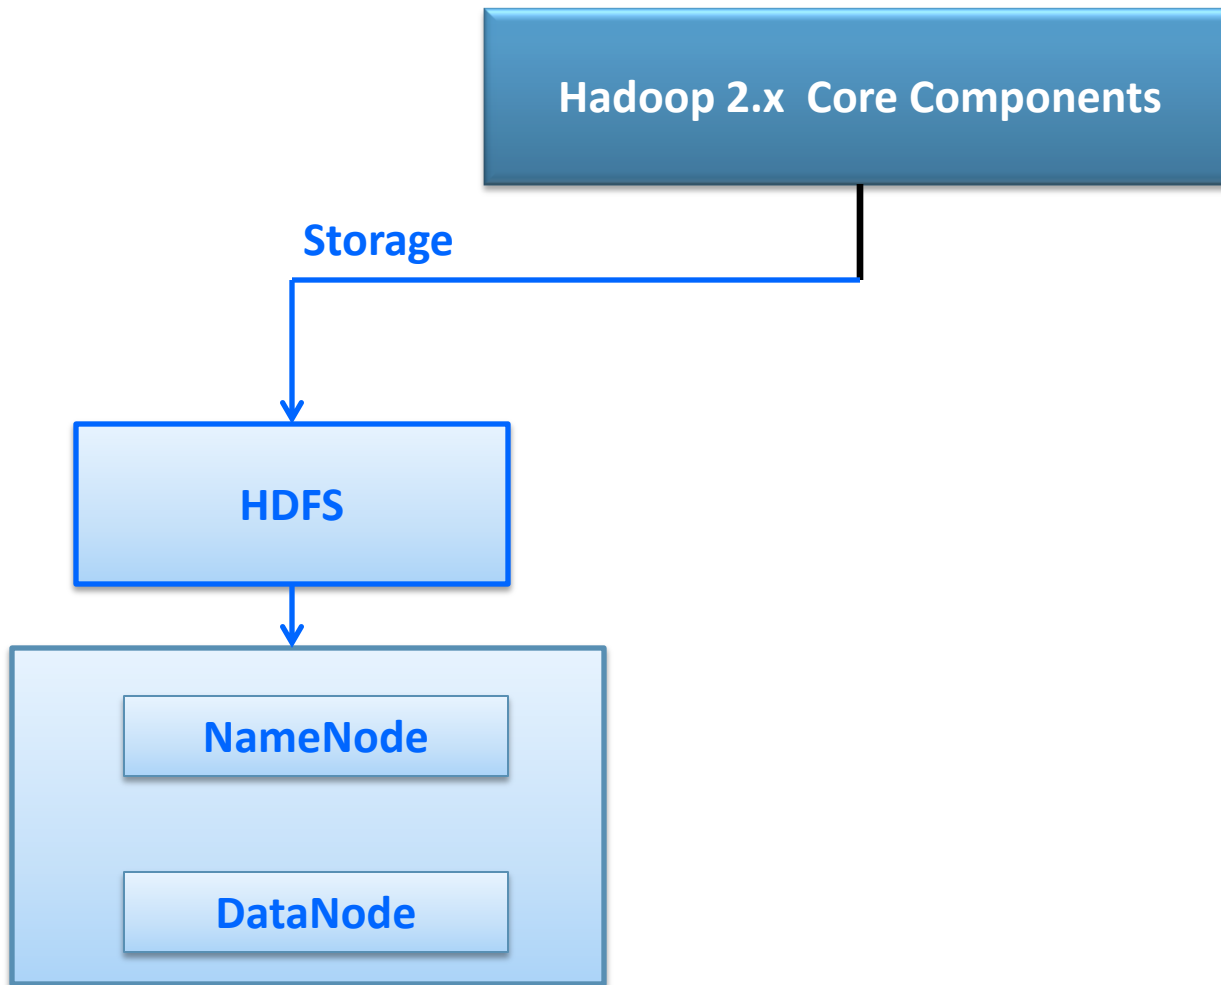


Core Components of Hadoop Cluster



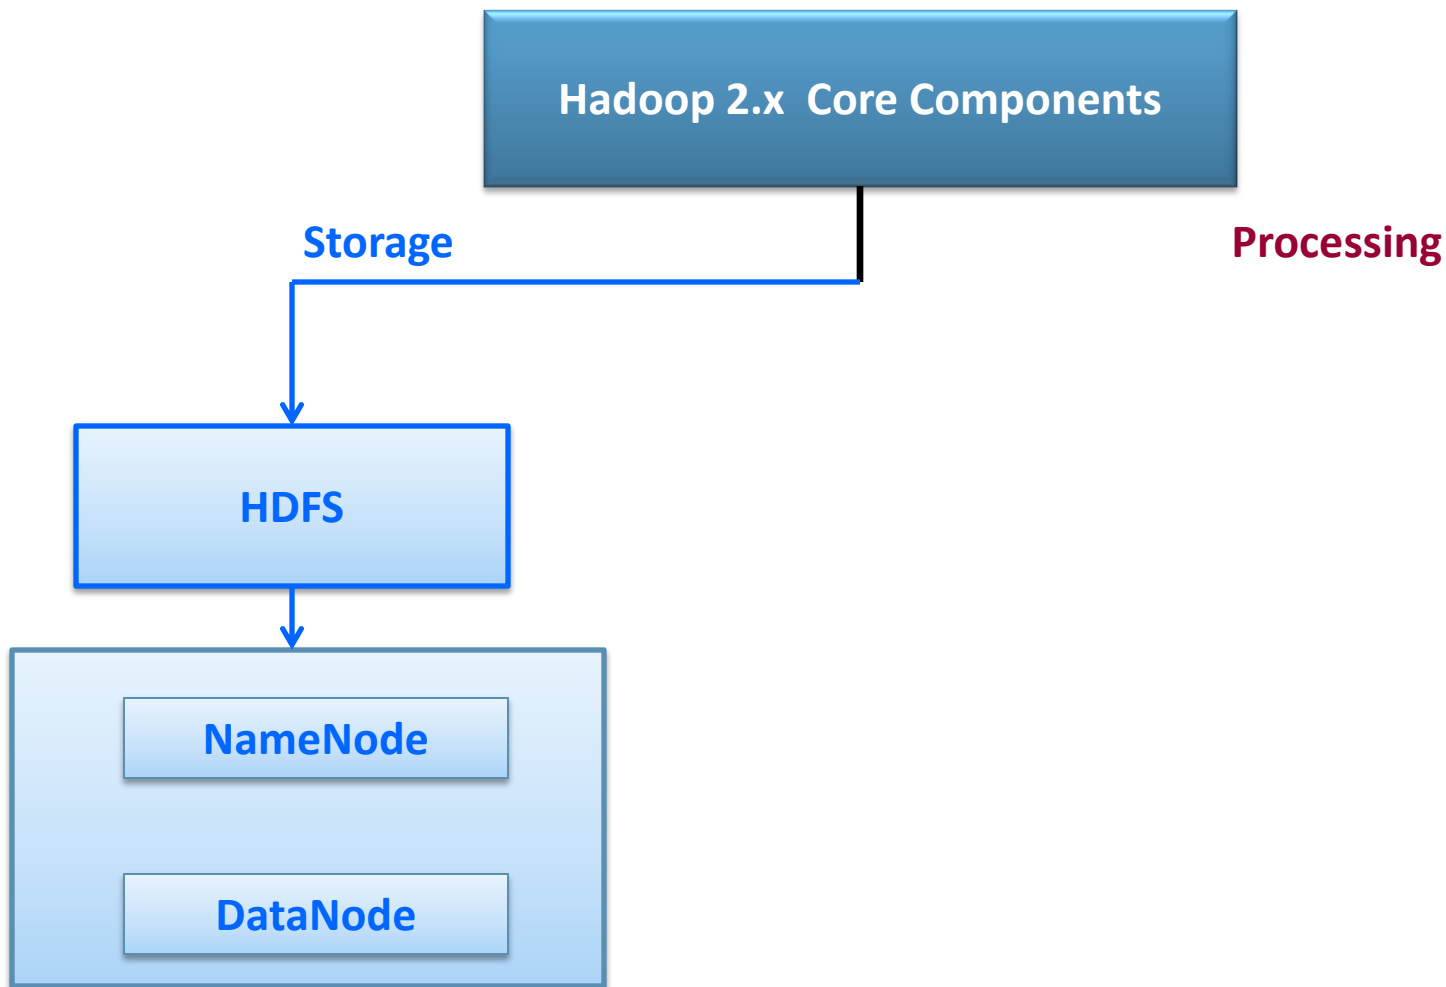


Core Components of Hadoop Cluster



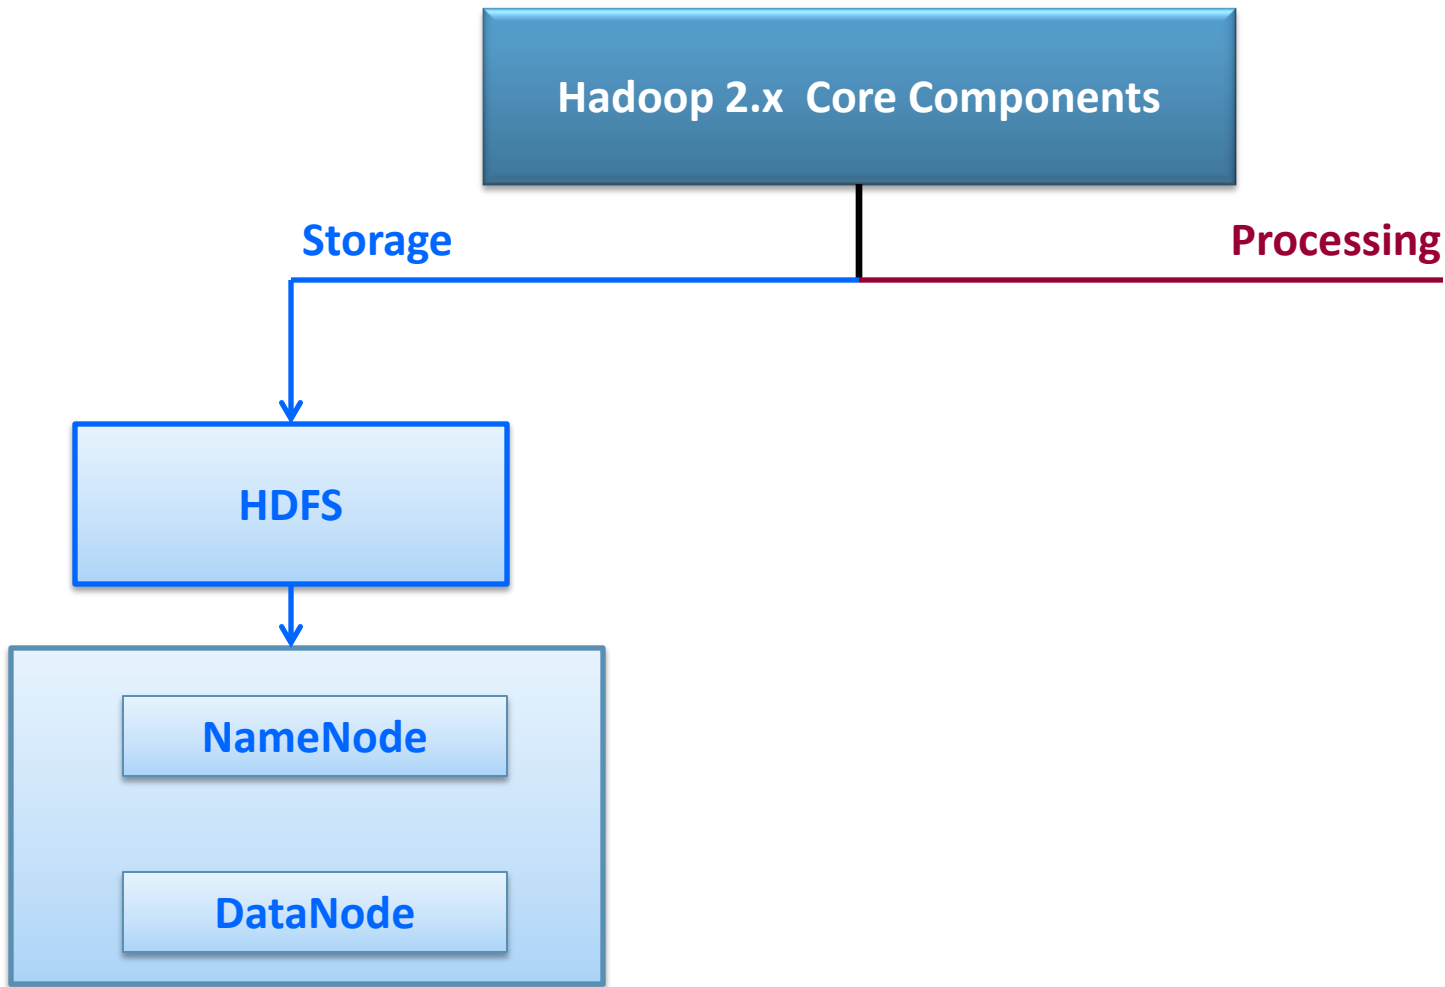


Core Components of Hadoop Cluster



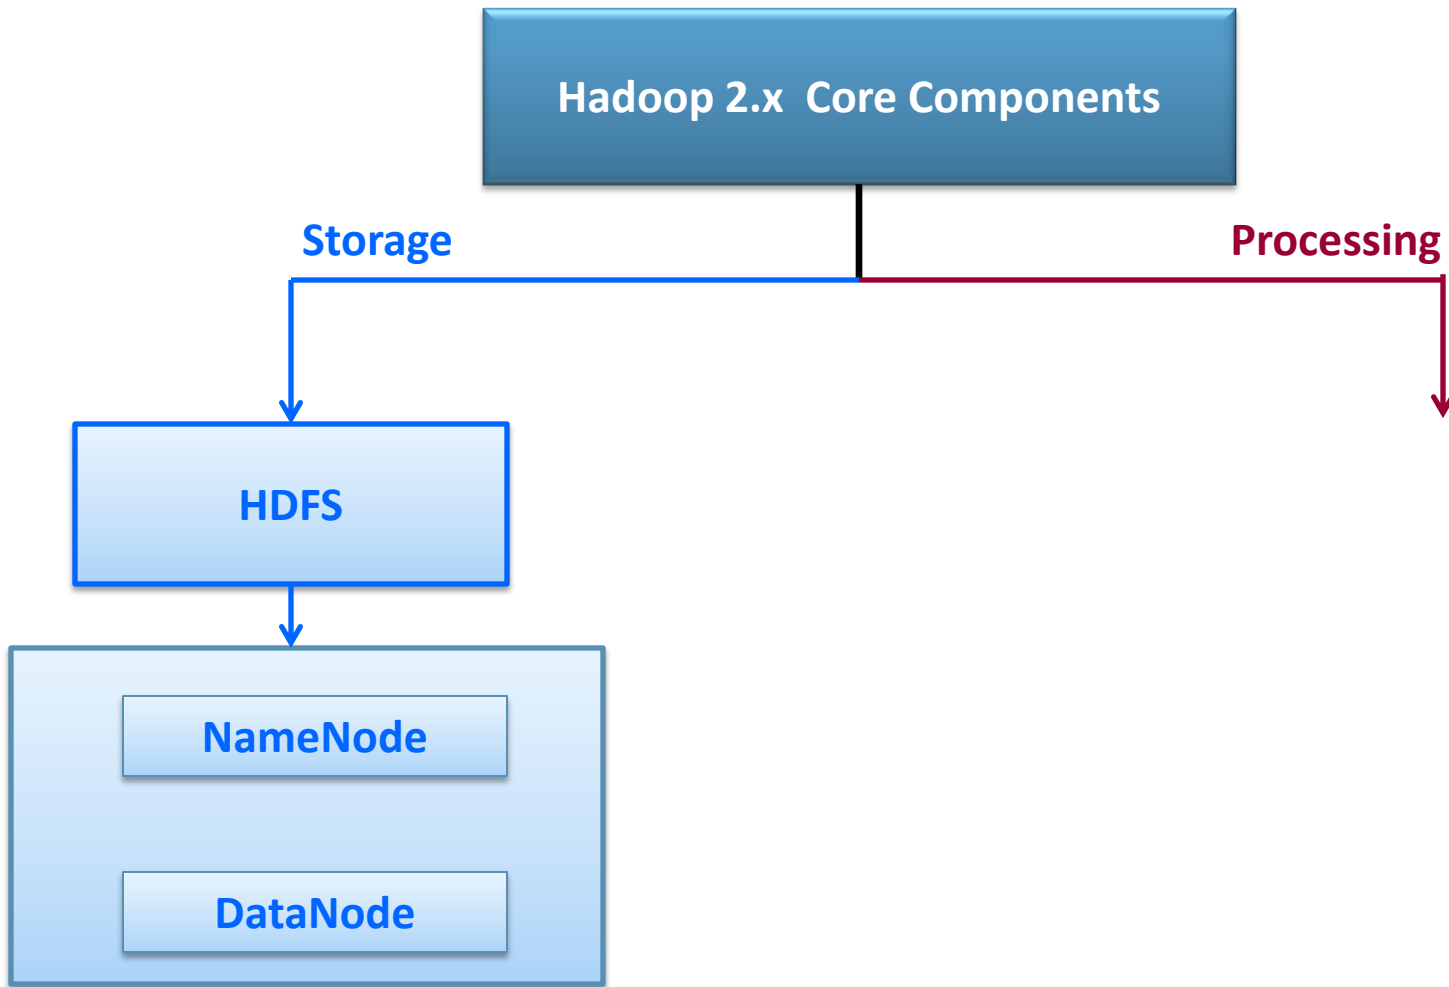


Core Components of Hadoop Cluster



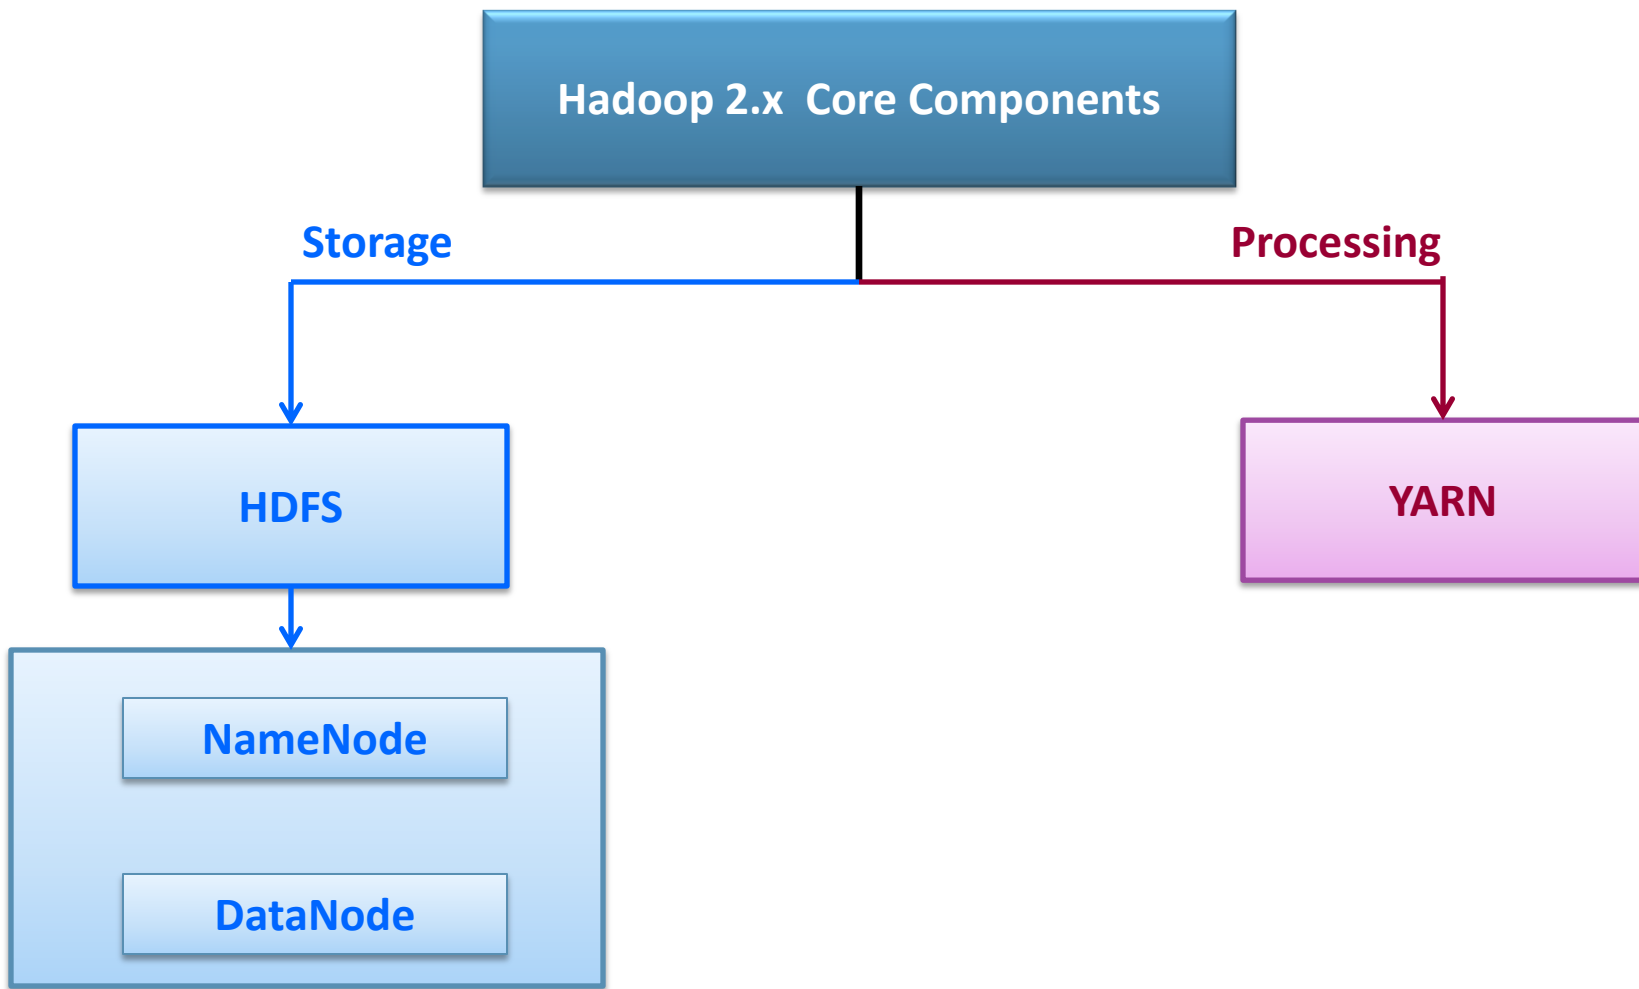


Core Components of Hadoop Cluster






Core Components of Hadoop Cluster



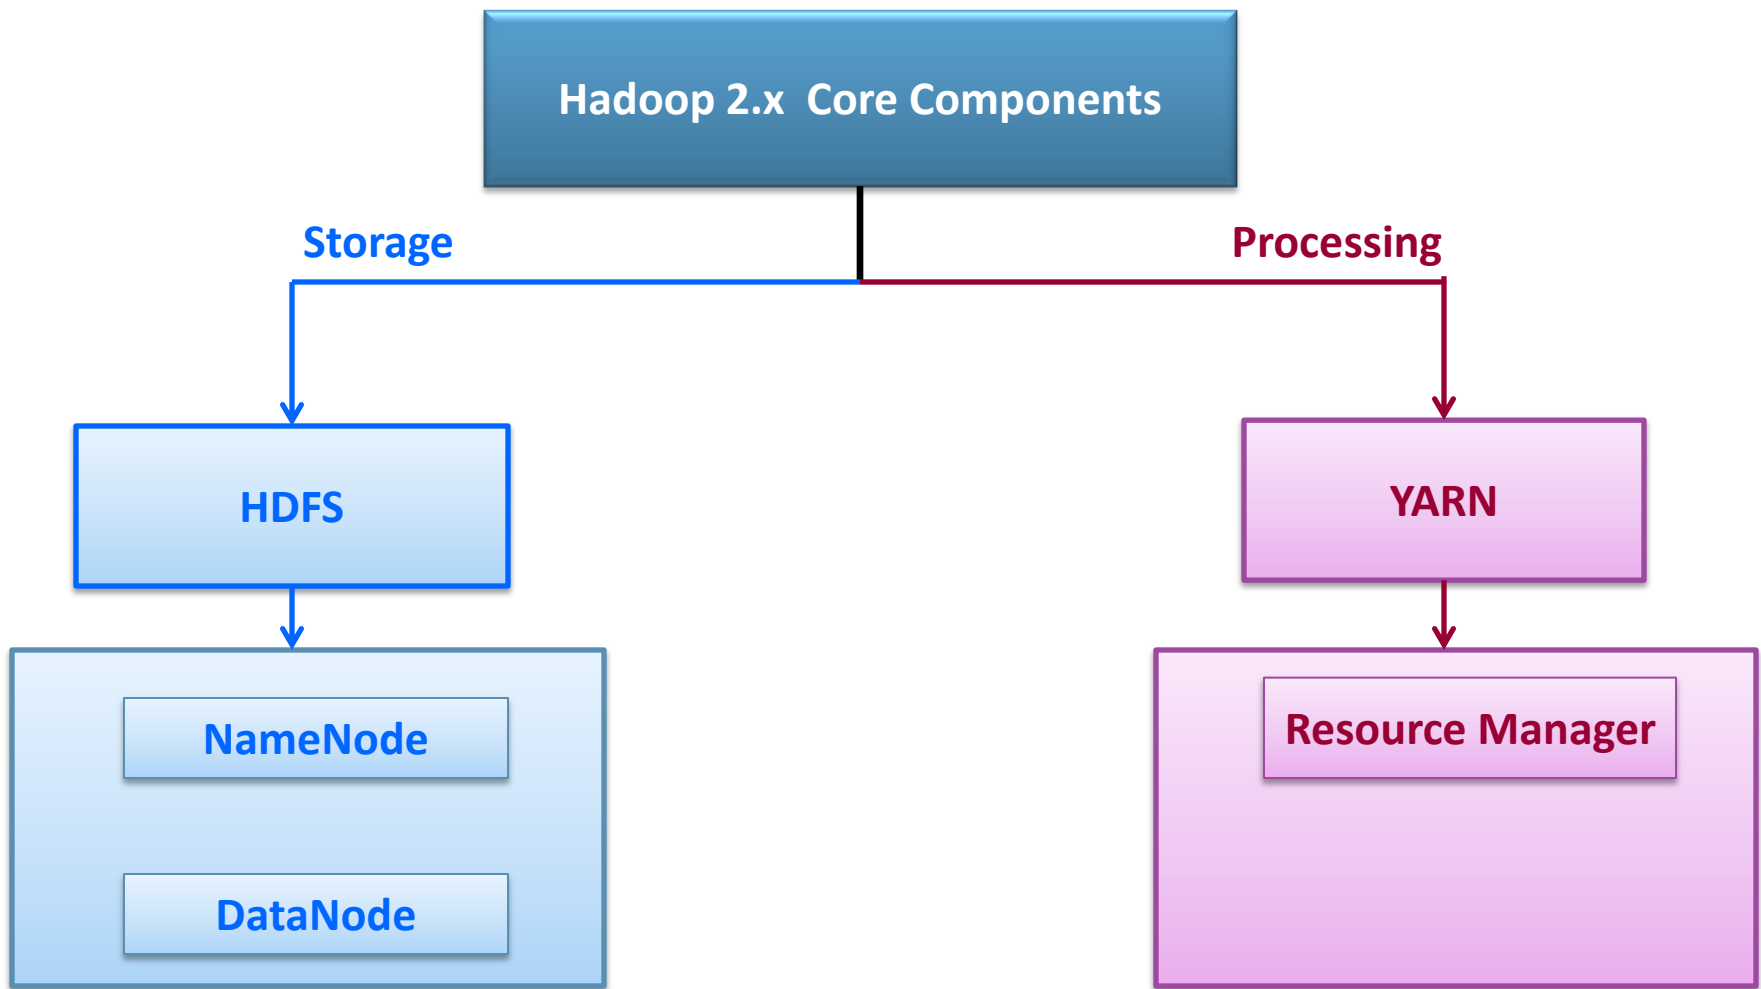


Core Components of Hadoop Cluster



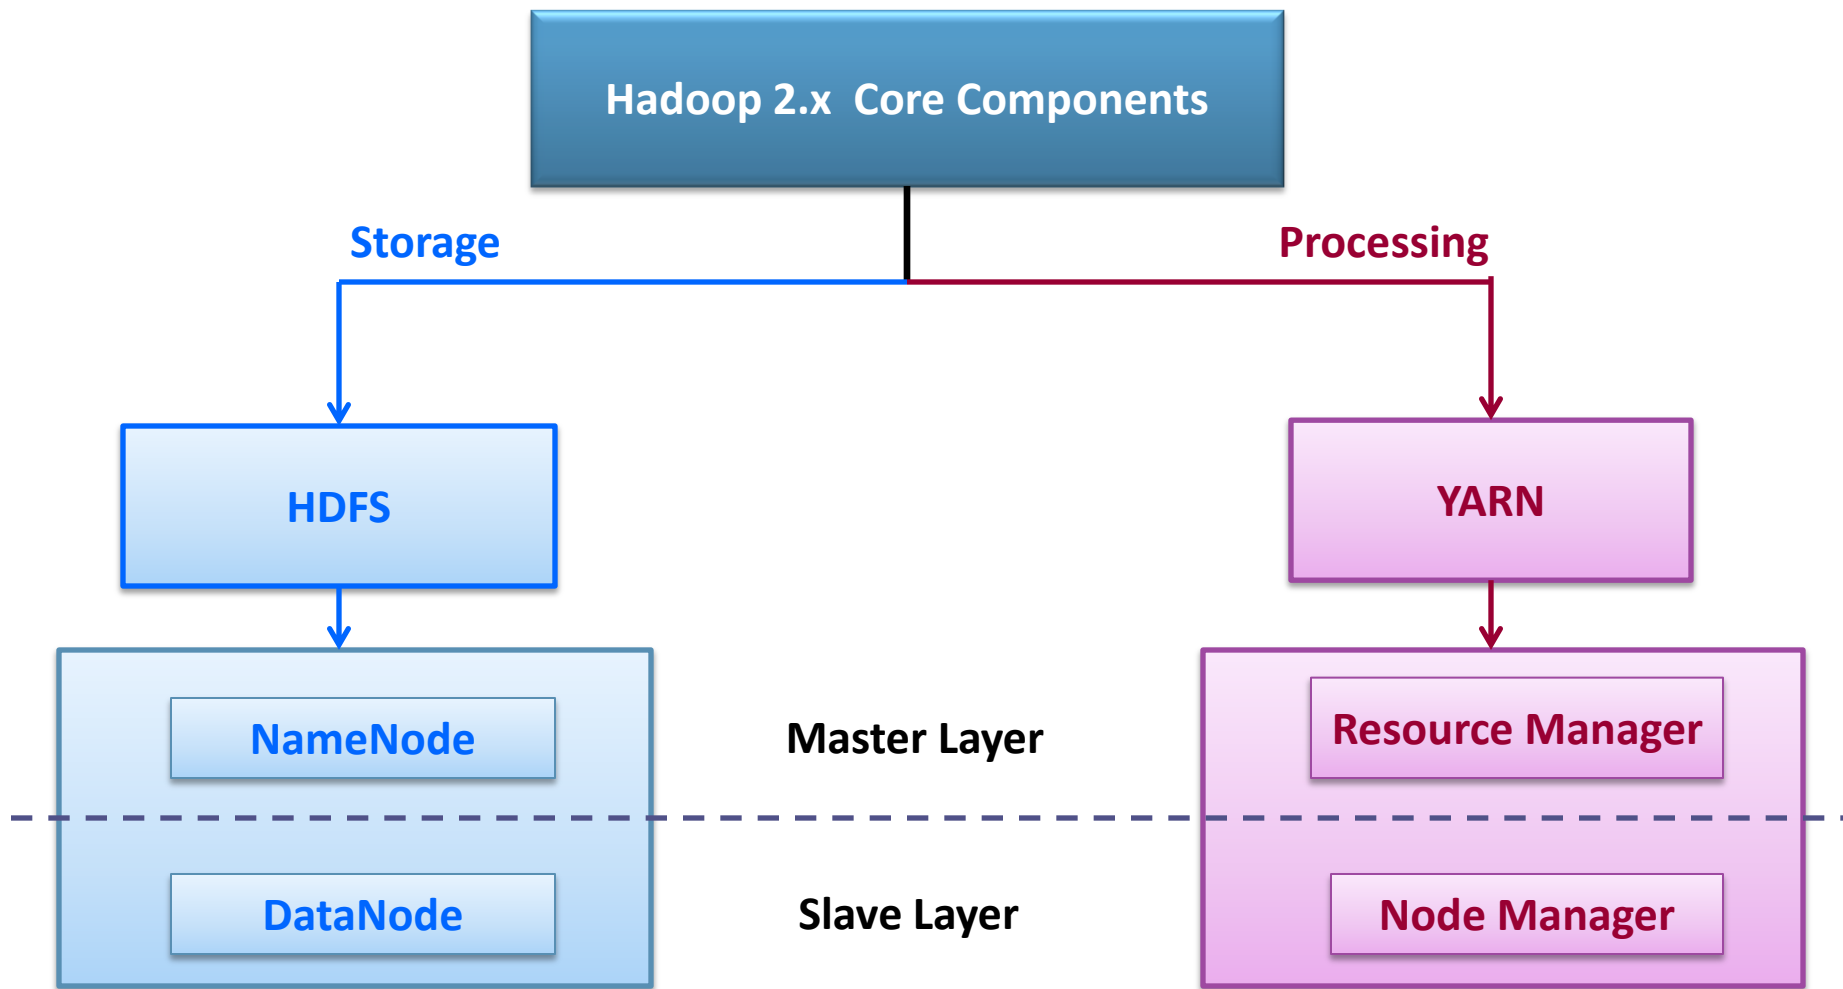


Core Components of Hadoop Cluster





Core Components of Hadoop Cluster



NameNode Backup in Hadoop1.x

- In Hadoop 1.x: Periodic updates of Edit Logs to Secondary NameNode are made:

All Namespace edits logged
as soon as an edit is done



Namespace logs
updated periodically

Active NameNode

Secondary
NameNode

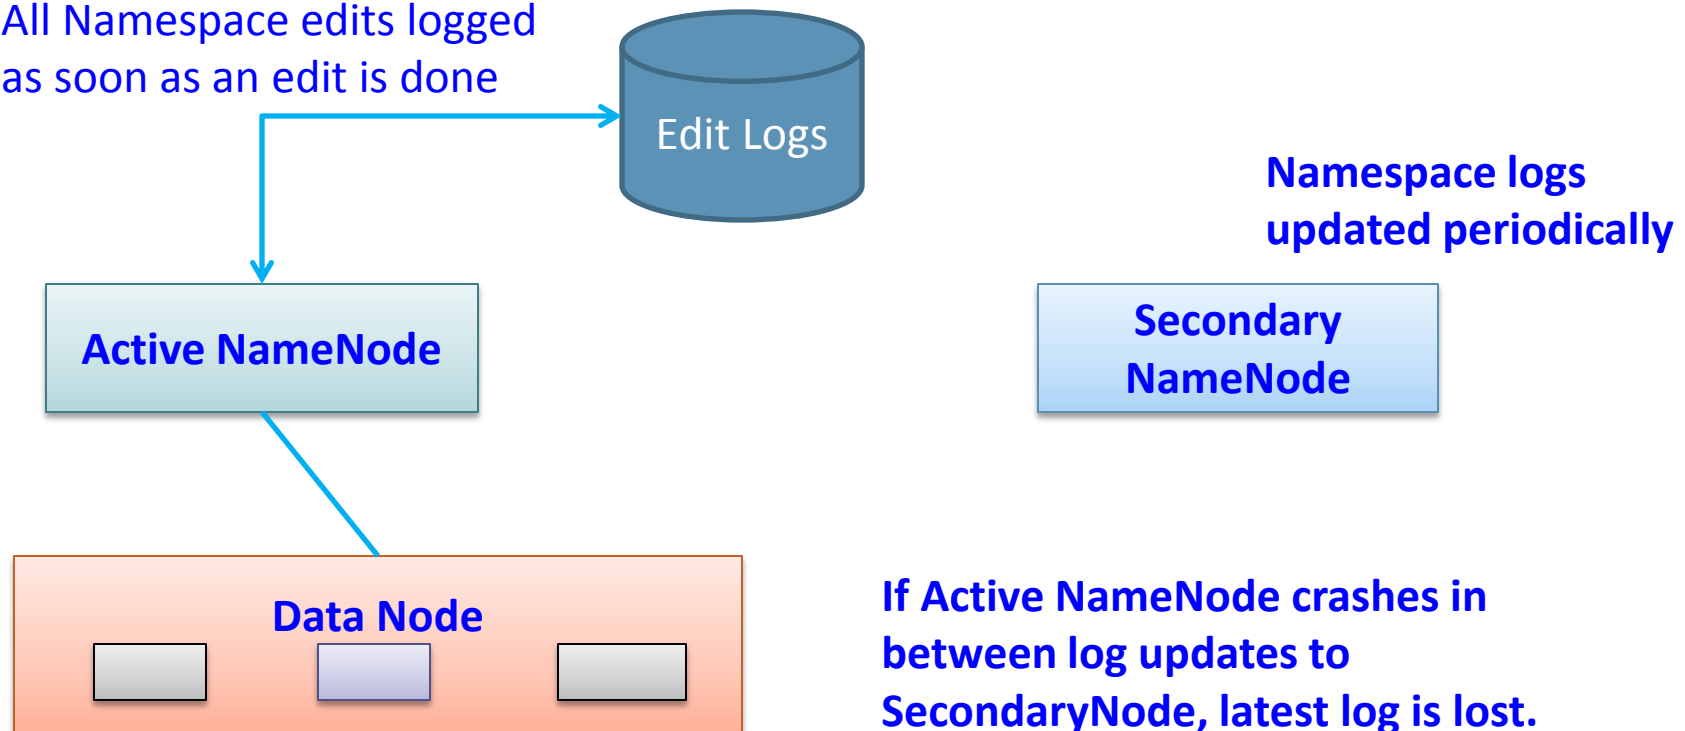
Data Node

If Active NameNode crashes in
between log updates to
SecondaryNode, latest log is lost.

NameNode Backup in Hadoop1.x

- In Hadoop 1.x: Periodic updates of Edit Logs to Secondary NameNode are made:

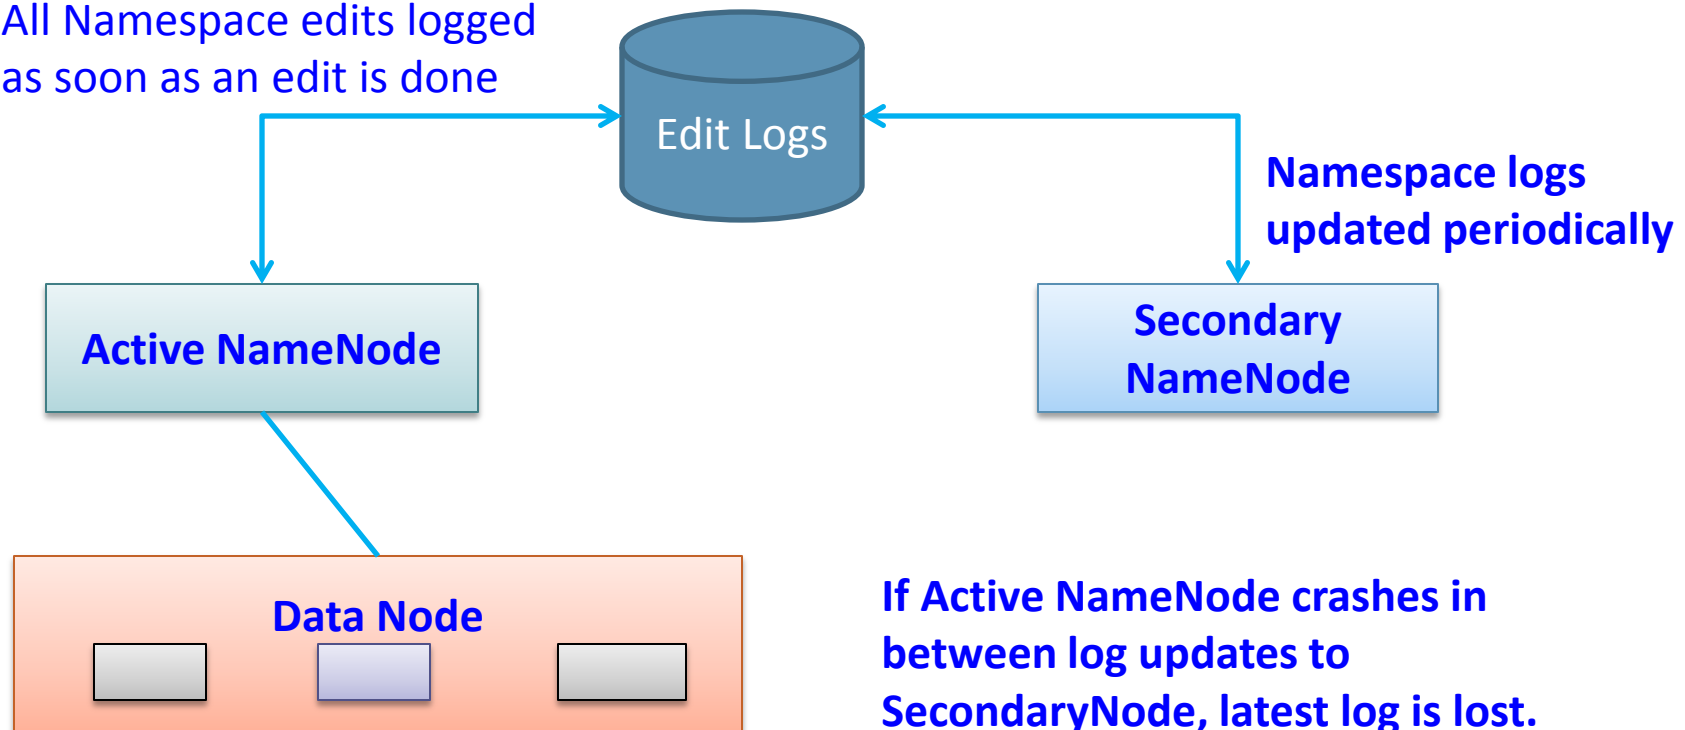
All Namespace edits logged
as soon as an edit is done



NameNode Backup in Hadoop1.x

- In Hadoop 1.x: Periodic updates of Edit Logs to Secondary NameNode are made:

All Namespace edits logged
as soon as an edit is done



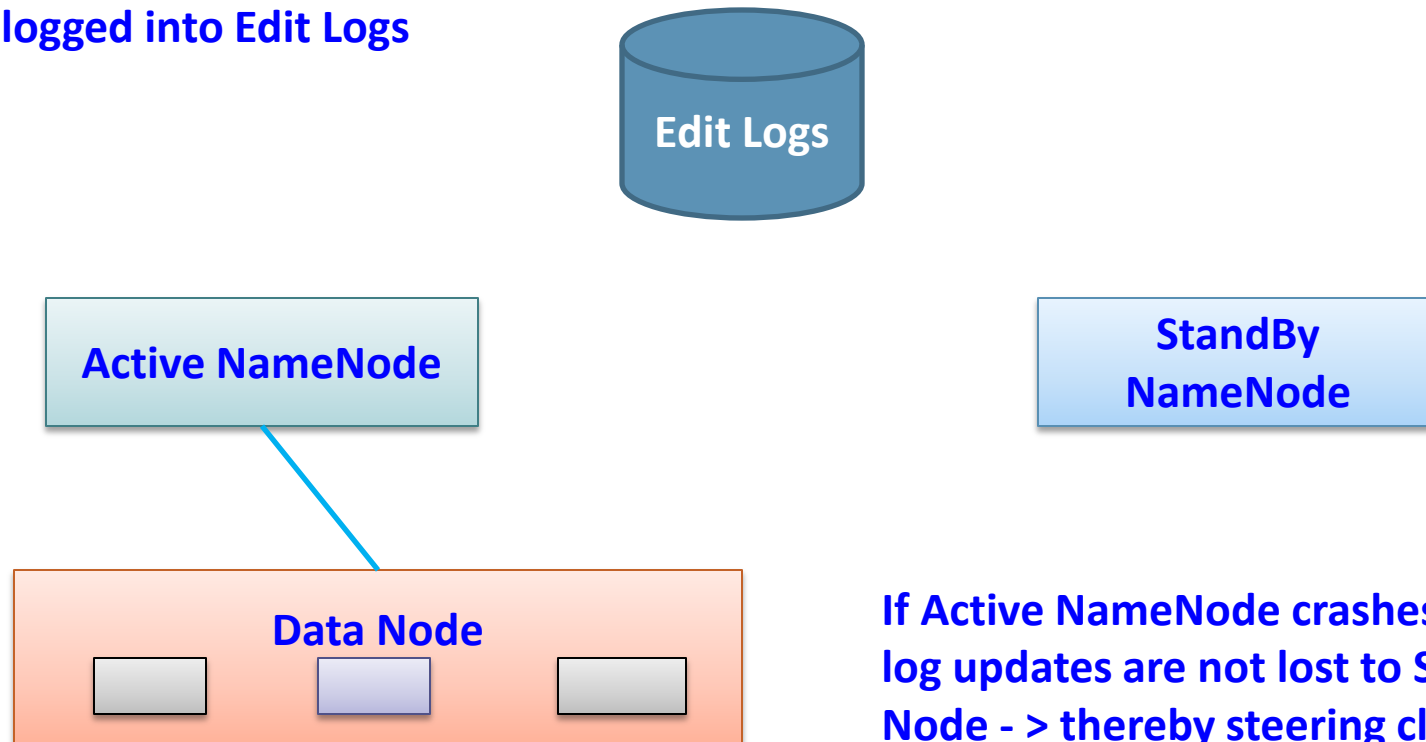
If Active NameNode crashes in
between log updates to
SecondaryNode, latest log is lost.



HDFS - High Availability Feature in Hadoop2.x

- StandBy NameNode in Hadoop 2.x:

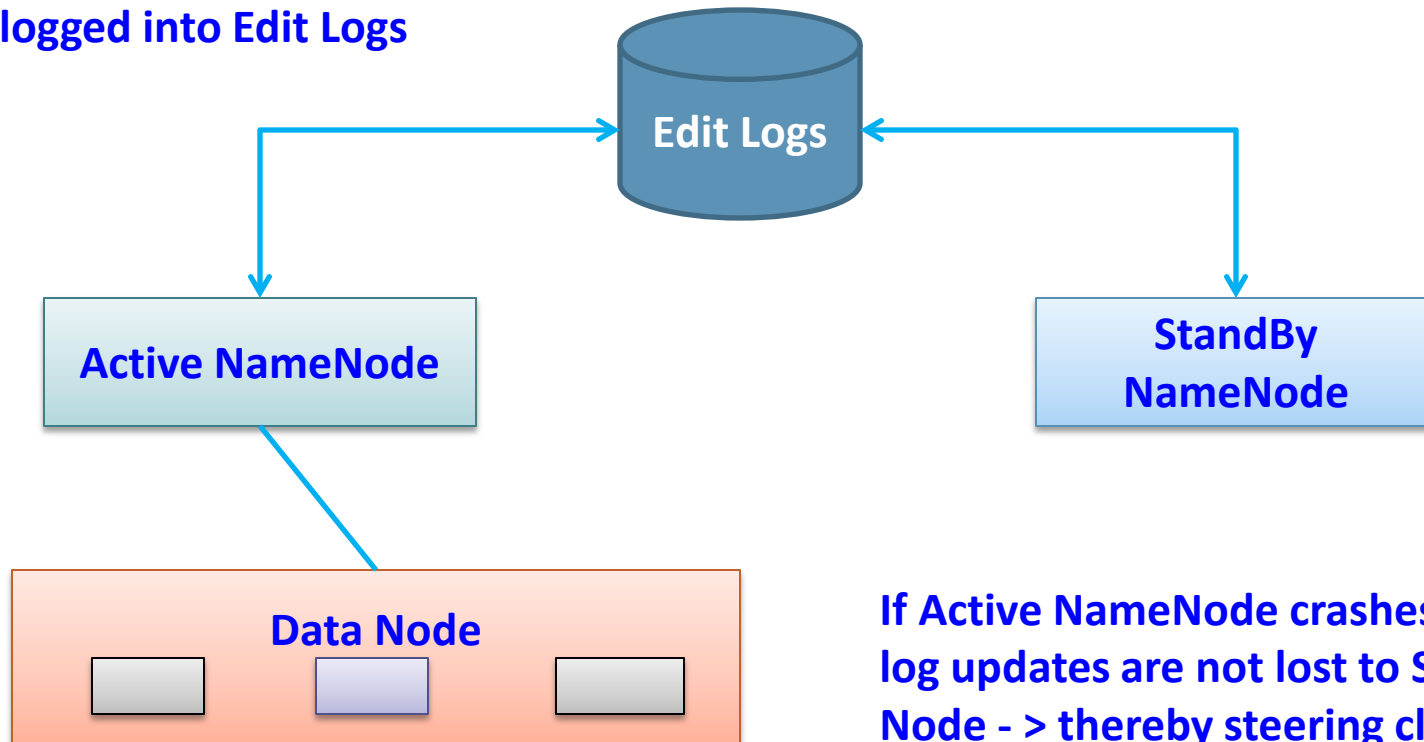
All Namespace edits
logged into Edit Logs



If Active NameNode crashes then log updates are not lost to StandBy Node - > thereby steering clear off - "NameNode single point of failure".

- StandBy NameNode in Hadoop 2.x:

All Namespace edits
logged into Edit Logs

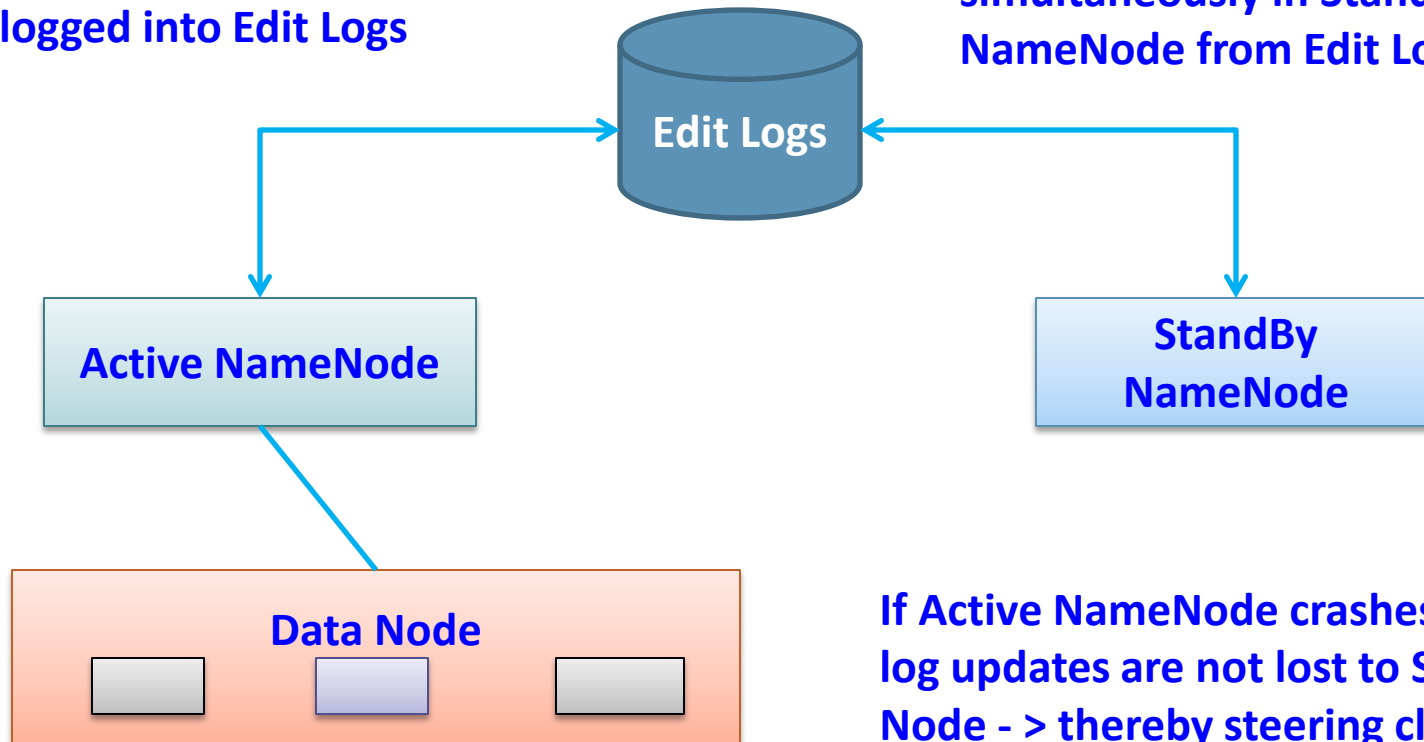


If Active NameNode crashes then log updates are not lost to StandBy Node - > thereby steering clear off - "NameNode single point of failure".

- StandBy NameNode in Hadoop 2.x:

All Namespace edits
logged into Edit Logs

Namespace logs (meta) updated
simultaneously in StandBy
NameNode from Edit Logs



If Active NameNode crashes then
log updates are not lost to StandBy
Node - > thereby steering clear off -
"NameNode single point of failure".




Introduction to a YARN Application

- In YARN a job is called an Application.
- For every job submitted to Resource Manager; it returns an acknowledgement in the form of an Application or Appld.
- YARN uses two daemons to launch & monitor containers:
 - Resource Manager
 - Node Manager
- A Container executes an application using constrained resources – CPU, memory etc.

ANATOMY of a YARN Application

Phase 1: Job Submission




Process of Job Submission:

1. Client calls the Submit() method on Job.
2. Job Submit() requests meta data from Resource Manager.
3. Resource Manager returns an Application Id.
4. Submit Application
5. Job Submit() gets the Input Splits from HDFS.
6. Job Submit() creates a directory in HDFS and copies into it the jar, Config files etc.



ANATOMY of a YARN Application

Phase 2: Job Initialization



Process of Job Initialization

1. Resource Manager finds a Node Manager and submits to it an Application() method to start the Application Master.
2. Node Manager launches Application Master in a Container.
3. Application Master initializes the job by creating a number of book keeping objects to track the job's progress.
4. Application Master retrieves Input Splits from HDFS.
5. Application Master now creates one map task per split & checks the mapreduce.job.reducers property and creates that many number of reducers.



Questions - Hadoop

- What is the approach to solve Big data problem?
- What are the core components of Hadoop?
- Do all nodes work all the time in Hadoop cluster?
- What is the role of Master Node in Hadoop Cluster?
- Can Slave node directly communicate with each other?
- How is the concept of High Availability more efficient than Secondary Name Node concept?



Next Session (3)

- HDFS Internals:
- Hadoop Distributed File System (HDFS):
 - Introduction
 - Design of HDFS
- Data Model:
 - HDFS Data Flow
 - Blocks in HDFS
 - HDFS high level architecture
 - Processing constraint with Hadoop physical block
 - Processing on Input Split.
 - Relation between Hadoop Block and Split
 - HDFS File-Write



Get In touch with AcadGild

Contact Info:

- **Website** : <http://www.acadgild.com>
- **LinkedIn** : <https://www.linkedin.com/company/acadgild>
- **Facebook** : <https://www.facebook.com/acadgild>
- **Support**: support@acadgild.com